

Optimization of LA in OLAP

2nd general debriefing meeting

Filipe Oliveira Sérgio Caldas

Advisors: Alberto Proença and José Nuno Oliveira

Table of Contents

- 1 Actual Progress
- 2 Translating TPC-H Query-1
- 3 Implementing and optimizing the LA operations
- 4 Dataset Characterization
- 5 Test Platform
- 6 Query-1 First Results
- 7 Next Steps

Actual Progress

Progress Until First Meeting

- LA Operations translated into our DSL
- Profiled the given dataset
- Our approach decisions
 - Use Sparse MKL Library
 - Use Quarks structure from Glib library
 - Use sparse matrices formats for matrices representations

Progress Until Now

- Sequential version of LA Operations implemented and optimized
- Querie-1 translated into our DSL
- Querie-1 First Results in our HPC test environment

Translating TPC-H Query-1

TPC-H Query-1 (simplified)

```
SELECT l_returnflag , l_linestatus , sum(l_quantity)
FROM LINEITEM_1
WHERE l_shipdate >= '1998-08-28' AND l_shipdate <= '1998-12-01'
GROUP BY l_returnflag , l_linestatus ;
```

Translates into LA Operations:

$$\underbrace{(L_{ReturnFlag} \nabla L_{LineStatus})}_{\text{projection}} \cdot \underbrace{[L]_{Shipdate}^{Shipdate \geq 1998-08-28} \cdot [L]_{Shipdate}^{Shipdate \leq 1998-12-01}}_{\text{selection}} \cdot \underbrace{[[L]]_{Quantity} \cdot !^{\circ}}_{\text{aggregation}}$$

Translating TPC-H Query-1

Translates into our DSL:

```
1 BEGIN
2 bitmap returnFlag, lineStatus, shipdate_gt, shipdate_lt;
3 matrix quantity;
4
5 returnFlag = tbl_read('lineitem.tbl', 8);
6 lineStatus = tbl_read('lineitem.tbl', 9);
7 quantity = tbl_read('lineitem.tbl', 4);
8 shipdate_gt = tbl_filter('lineitem.tbl', 10, >=, '1998-08-28' );
9 shipdate_lt = tbl_filter('lineitem.tbl', 10, <=, '1998-12-01' );
10
11
12 bitmap projection, selection;
13 matrix aggregation, intermediate_result, final_result;
14
15 vector bng;
16 bng = bang(6);
17
18 selection = shipdate_gt * shipdate_lt;
19 projection = returnFlag krao lineStatus;
20 aggregation = quantity * bng;
21 intermediate_result = projection * selection;
22 final_result = itermediate_result * aggregation;
23
24 tbl_write(final_result, 'final_result.tbl');
25 END
```

Implementing and optimizing the LA operations

- Several LA approaches were considered
- La operations required: dot product, Kronecker Product, Hadamard Product, Khatri-Rao Product
- Based on the special matrices property (**1 element per column**) all 4 products were simplified (**not enough to auto-vectorize**)
 - All produced matrices have to be squared (if not pad with 0s)
- Opportunity to use SSE/AVX vector instructions in all 3 products
 - Enforced vectorization of loops
 - All Data is aligned during creation accordingly to cache line size in order to assist vectorization (**Allocation alignment**)
 - Instructed the compiler to assume that all CSR arrays are aligned on an 32-byte boundary (**Access alignment**)

Implementing and optimizing the LA operations

- Khatri-Rao (used in query-1) optimization report:

```
94 Begin optimization report for: csr_krao(float *, int *, int *, int, int, int, float *, int *, int *, int, int,
95
96     Report from: Vector optimizations [vec]
97
98
99 LOOP BEGIN at olap_search.c(538,13)
100 remark #15388: vectorization support: reference C_IA1 has aligned access
101 remark #15388: vectorization support: reference A_IA1 has aligned access
102 remark #15388: vectorization support: reference C_csc_values has aligned access [ olap_search.c(540,20) ]
103 remark #15388: vectorization support: reference B_csc_values has aligned access [ olap_search.c(540,20) ]
104 remark #15388: vectorization support: reference A_csc_values has aligned access [ olap_search.c(540,20) ]
105 remark #15388: vectorization support: reference C_JA1 has aligned access [ olap_search.c(541,13) ]
106 remark #15388: vectorization support: reference B_JA1 has aligned access [ olap_search.c(541,13) ]
107 remark #15388: vectorization support: reference A_JA1 has aligned access [ olap_search.c(541,13) ]
108 remark #15305: vectorization support: vector length 8
109 remark #15309: vectorization support: normalized vectorization overhead 0.200
110 remark #15301: FUSED LOOP WAS VECTORIZED
111 remark #15448: unmasked aligned unit stride loads: 5.
112 remark #15449: unmasked aligned unit stride stores: 3.
113 remark #15475: --- begin vector loop cost summary ---
114 remark #15476: scalar loop cost: 18.
115 remark #15477: vector loop cost: 3.120.
116 remark #15478: estimated potential speedup: 5.380.
117 remark #15488: --- end vector loop cost summary ---
118 LOOP END
```

Dataset Characterization

TPC-H Benchmark Tables

- Used in TPC-H Query-1:
 - **lineitem.tbl**
- Other TPC-H Benchmark Tables:
 - customer.tbl
 - nation.tbl
 - part.tbl
 - region.tbl
 - orders.tbl
 - partsupp.tbl
 - supplier.tbl

Datasets

- Generated data for different sizes: 1GB, 2GB, 4GB, 8GB, 16GB and 32 GB, producing 6 distinct LINEITEM DB tables.
- To speed up queries, indexes were created for all DB Tables:
 - f.e.: `CREATE INDEX l_orderkey_idx8 on LINEITEM_8 (l_orderkey);`

Environment Test Platform (Hardware)

System	compute-652-1
# CPUs	2
CPU	Intel® Xeon® E5-2670v2
Architecture	Ivy Bridge
# Cores per CPU	10
# Threads per CPU	20
Clock Freq.	2.5 GHz
L1 Cache	10 x 32 KB instruction caches 10 x 32 KB data caches
L2 Cache	2,560 KB 256 KB per core
L3 Cache	25 MB
Inst. Set Ext.	AVX, SIMD
#Memory Channels	4
Vendors Announced Peak Memory BW	59.7 GB/s
Measured Peak Memory BW	57.18 GB/s

Environment Test Platform (Software)

- LA OLAP:

- Compiler: ICC version 16.0.0 (GCC version 4.4.6 compatibility)
 - no vectorization: `-O3 -std=c99 -no-vec -farray-notation`
 - vectorization: `-O3 -std=c99 -farray-notation -xAVX -vec-report7`
- Intel® MKL Version 11.3
 - Link line: `-lmkl_intel_lp64 -lmkl_core -lmkl_sequential -lpthread -lm`

- PostgreSQL: version 9.6+

- Built with the following dependencies:
 - GCC version 4.9.0
 - Python 2.6.6
- Why PostgreSQL 9.6+ ?
 - Open-Source
 - PostgreSQL has Parallel-Query available¹

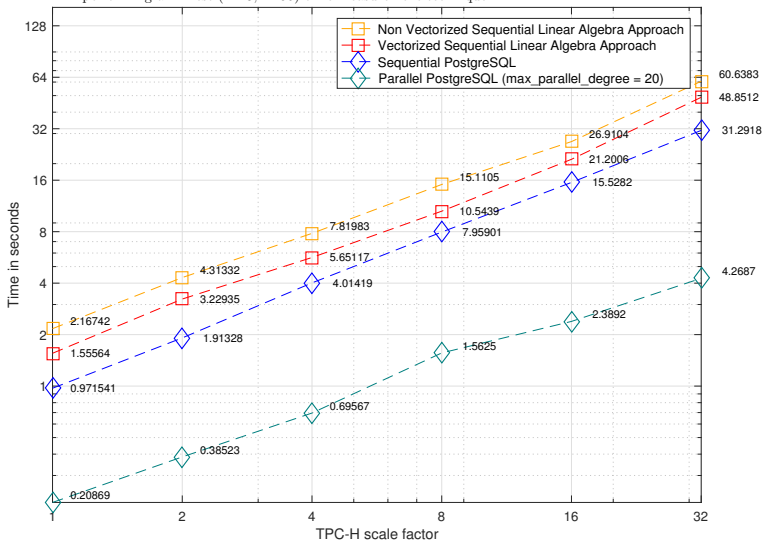
- Benchmarking:

- 22 TPC-H Benchmark queries

¹ <http://www.postgresql.org/docs/9.6/static/runtime-config-resource.html>

Query-1 First Results in our HPC test environment

TPC-H benchmark simplified query-1 time for solution analysis
for different scale factors, between Linear Algebra approach vs Relational Algebra approach,
performing a K-Best (K=3,N=50) time measurement technique



Planning our next steps

- Multithreaded parallelism
 - Expected high speedup values
- Improve data locality
 - Implement all products in BSR format
 - Irregular access patterns decreases performance
- Analyse Gather/Scatter analysis for distributed memory parallelism
 - Detect possible latency or bandwidth issues

Optimization of LA in OLAP

2nd general debriefing meeting

Filipe Oliveira Sérgio Caldas

Advisors: Alberto Proença and José Nuno Oliveira